UNIVERSITÀ
DEGLI STUDI
DI PADOVA

N-Body Simulation: Design, Optimizations and Benchmarking

M. Sc. Physics of Data, University of Padova

Alessio

July 23, 2025

Outline

1. Introduction and Motivation
2. Project Evolution
3. Memory Management and Optimizations
4. Force and Integration Methods
5. Limitations and Lessons Learned
6. Conclusions

Introduction and Motivation

- Started as a simple single-file C++ code for Voyager II trajectory
- Evolved into a modular, configurable N-body simulator
- Main goals:
  - Flexible and extensible framework
  - Fair CPU vs GPU performance comparison
  - Advanced optimizations and interactive visualization

Project Evolution

- Initial single-file C++ implementation
- Output separation and modularization
- Integration of OpenMP, CUDA, (AVX2)
- Refactoring, CMake build system, JSON configuration
- Interactive visualization with OpenGL

General Architecture

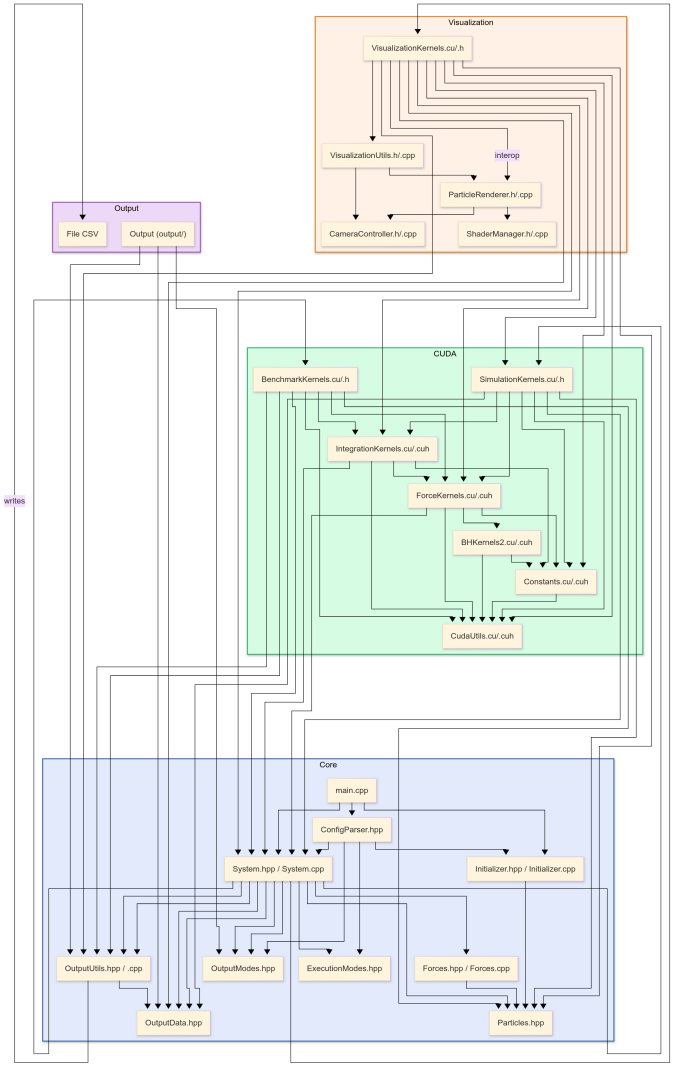

- **Core**: simulation logic, data structures, and CPU part
- **CUDA**: GPU kernels, utilities and host functions
- **Visualization**: OpenGL and OpenGL/CUDA interop
- **Output**: CSV output, dynamical buffer to use for other integrations
- **Configuration**: single JSON file

Configuration and Initialization

- Single JSON configuration file
- Flexible initialization: random, file-based, or custom implementations (galaxy, stellar systems, etc.)

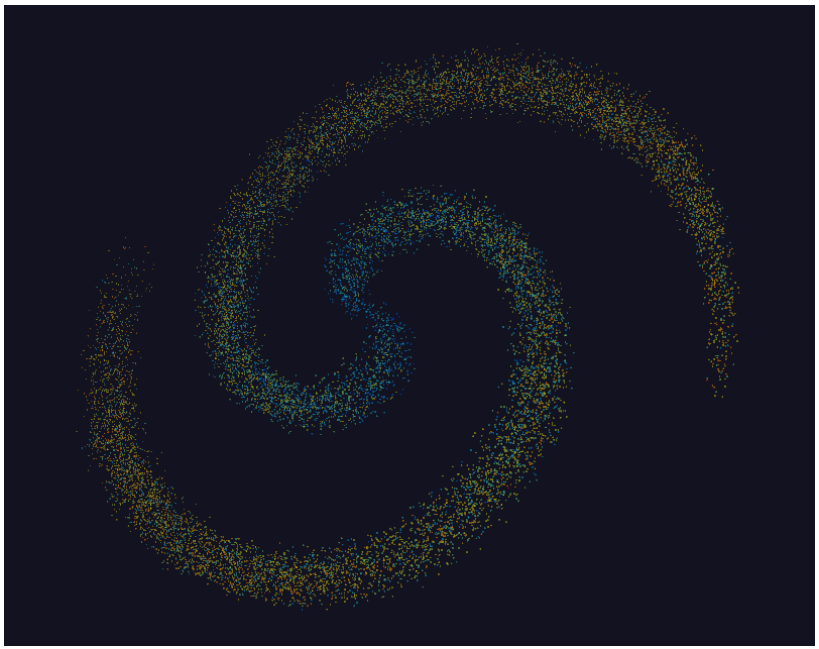

Example of a galaxy configuration

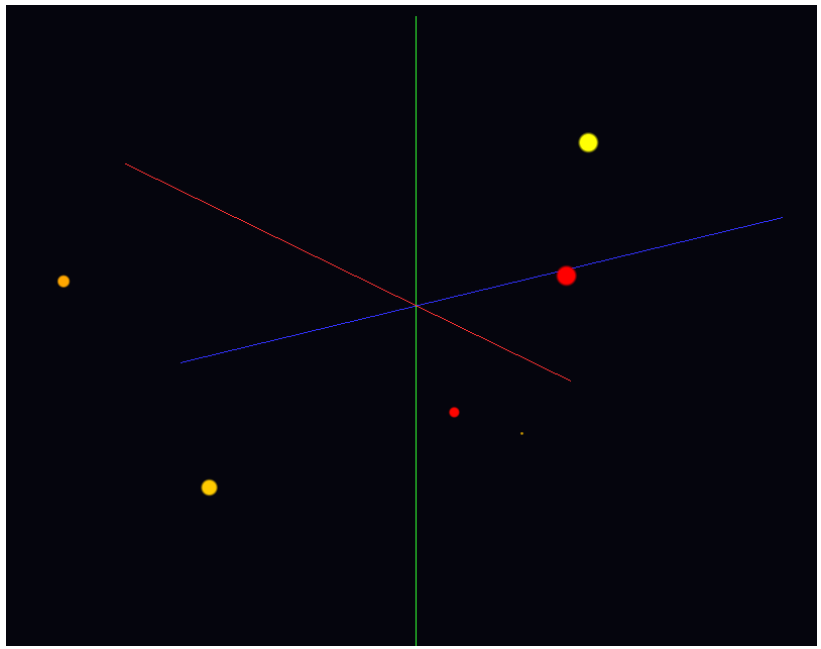

Example of random initialization

Memory Management and Optimizations

- Structure of Arrays (SoA) for efficient access and fully coalesced memory with double4
- Pinned memory and CUDA streams
- AVX2 vectorization and OpenMP threading
- Low-ish memory footprint: 1.1 GiB for 1B particles
- In CSV output mode, System is copied only at the start $\rightarrow$ multiple cuda streams for memory and compute management.
- In visualization mode, System is copied every frame to a dynamic buffer, which is then used for rendering $\rightarrow$ worst case scenario $\Rightarrow$ OpenGL-cuda interop and separate rendering stream $\rightarrow$ less bottleneck.

Force and Integration Methods

- Pairwise, Adaptive Mutual Softening, Barnes-Hut (no GPU implementation)
- Configurable physical constants and softening
- Multiple integration schemes: Euler (not symplectic), Velocity Verlet (symplectic)

Pairwise Newtonian:

$$\vec{a}_i = G \sum_{j \neq i} m_j \frac{\vec{r}_j - \vec{r}_i}{|\vec{r}_j - \vec{r}_i|^3}$$

Adaptive Mutual Softening:

$$\vec{a}_i = G \sum_{j \neq i} m_j \frac{\vec{r}_j - \vec{r}_i}{\left(|\vec{r}_j - \vec{r}_i|^2 + \epsilon_{ij}^2\right)^{3/2}}$$

$$\epsilon_{ij} = \max\left(\eta \, |\vec{r}_j - \vec{r}_i| \left(\frac{m_i + m_j}{3\langle m \rangle}\right)^{1/3}, \epsilon_{\min}\right)$$

where G is the gravitational constant, ϵ_{ij} is the adaptive softening, η and $\epsilon_{\min}$ are tunable parameters, and $\langle m \rangle$ is the average mass.

Block Size Auto-Tuning

- **block size selection** is crucial for GPU efficiency in these problems.
- Two strategies compared:
  - Occupancy API (NVIDIA)
  - Heuristic (custom, optimized for this workload)

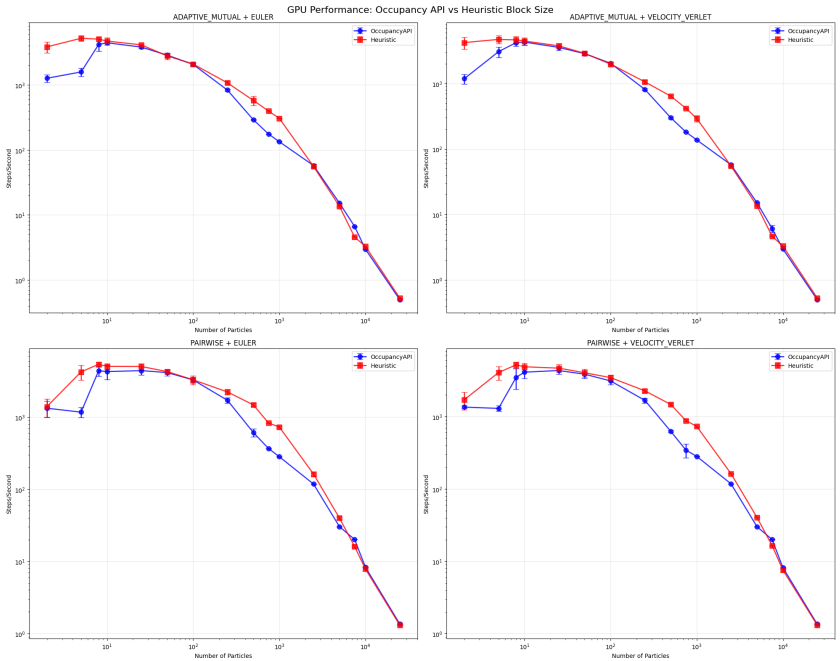

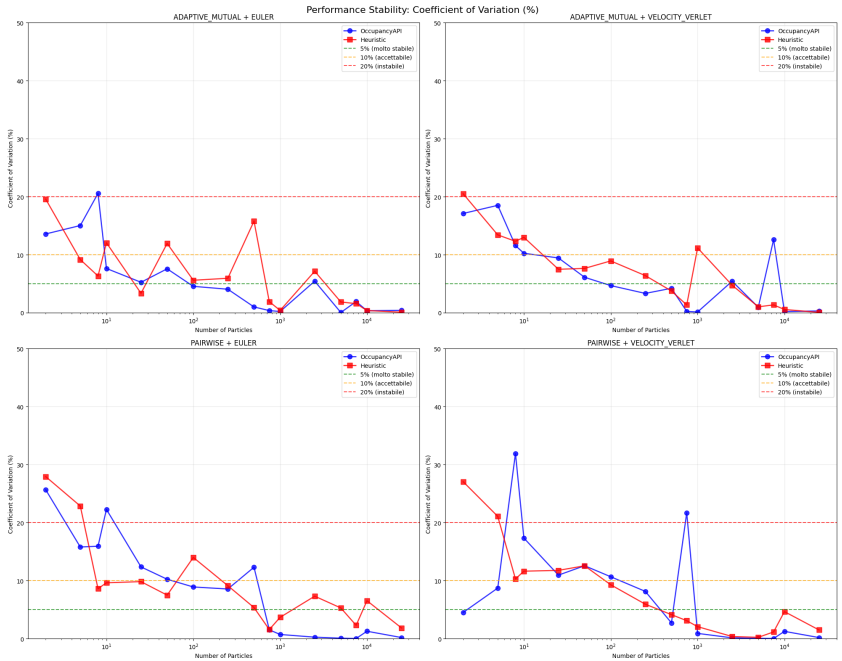

Block Size Auto-Tuning

Grid size:

$$\text{gridSize} = \left\lceil \frac{n}{\text{blockSize}} \right\rceil$$

The heuristic is robust, saturates SMs, and respects shared memory limits. The grid size formula is exact for 1D N-body kernels.

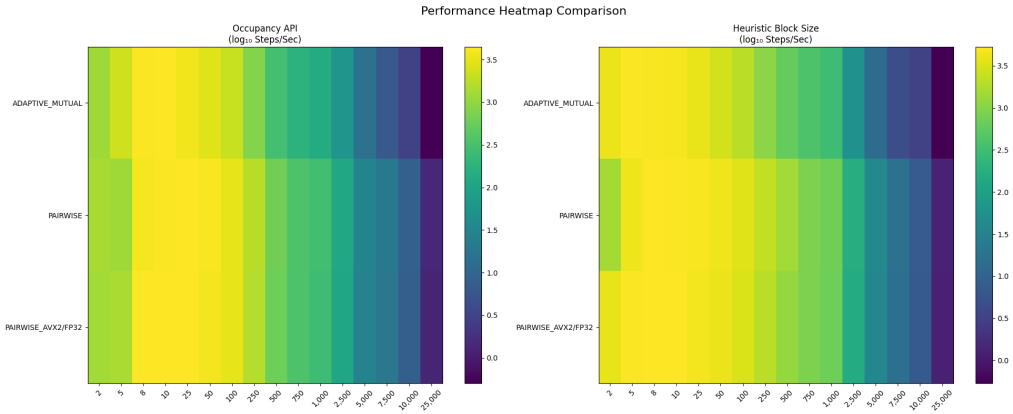

Heuristic Block Size Selection (pseudocode)

```
blockSize = 1024
if n ≤ 32:  return 32
while sharedMem(blockSize)
> maxShared && blockSize
> warp:
    blockSize //= 2
blockSize = min(blockSize,
maxThreads)
while numBlocks(blockSize)
< 2*numSMs && blockSize
> warp:
    blockSize //= 2
blockSize = max(blockSize,
warp)
return blockSize
```

Basic CPU-GPU Comparison Benchmark

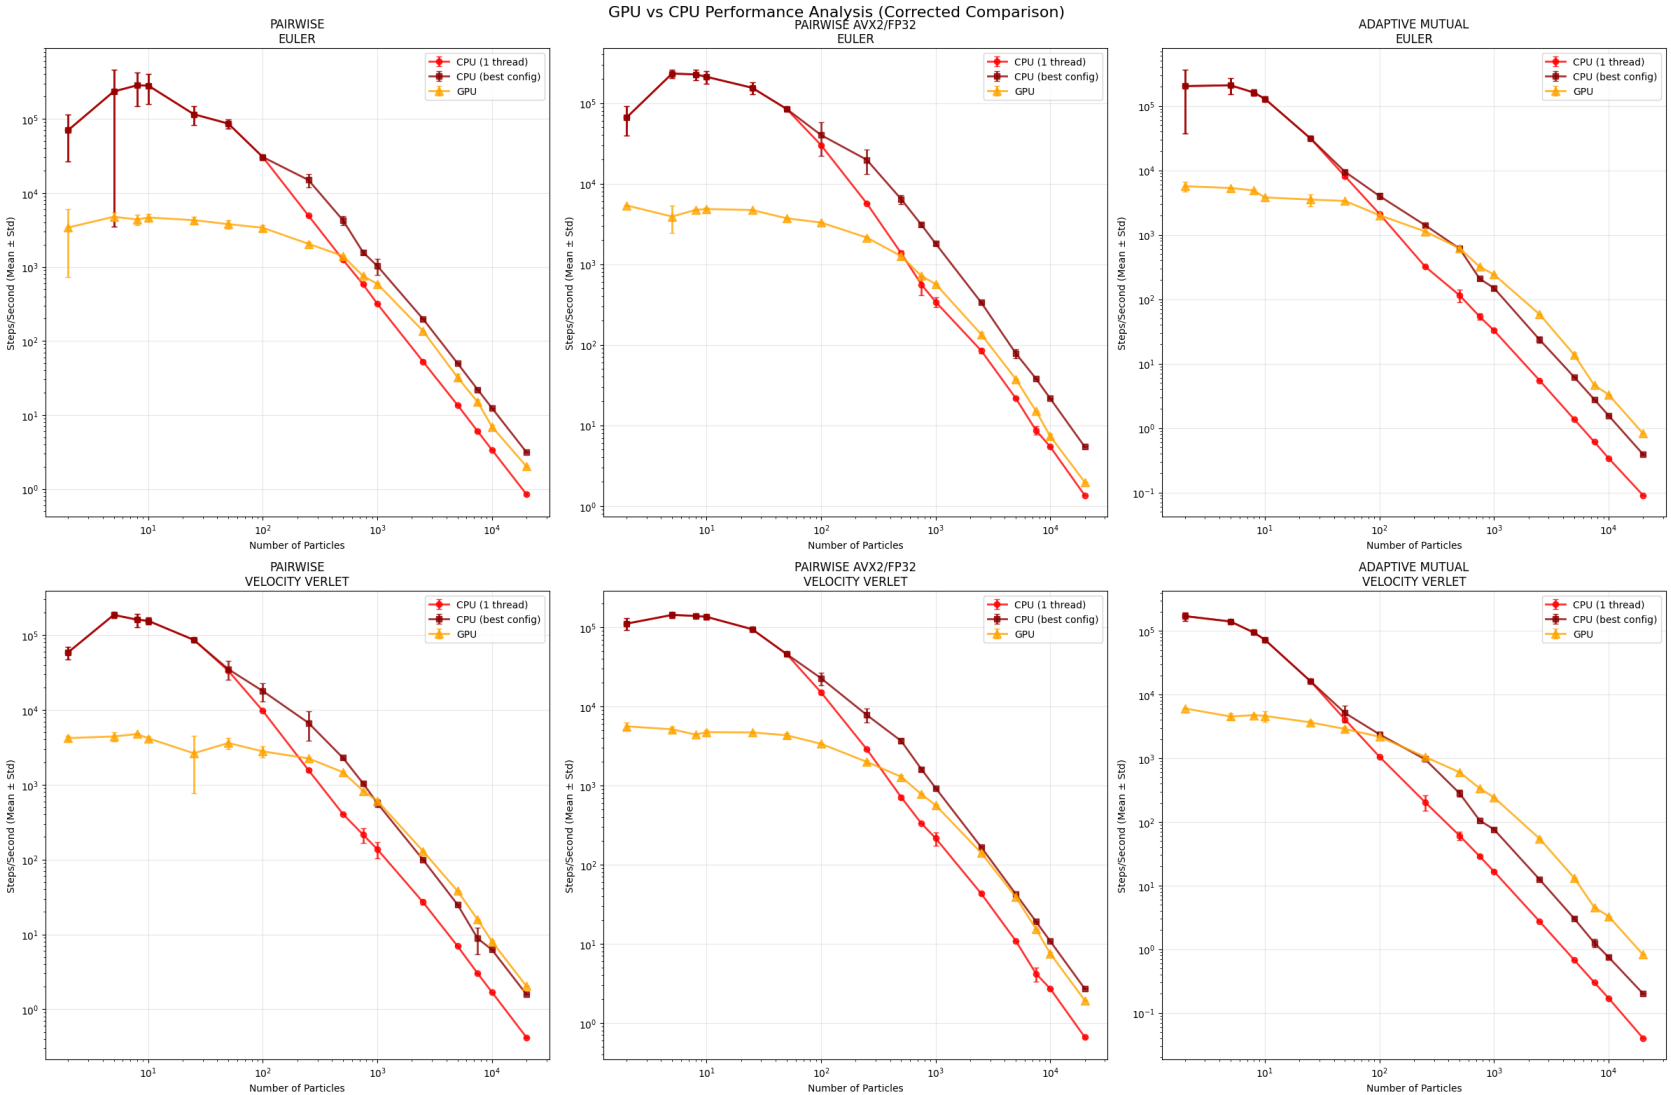

Basic CPU-GPU Comparison Benchmark

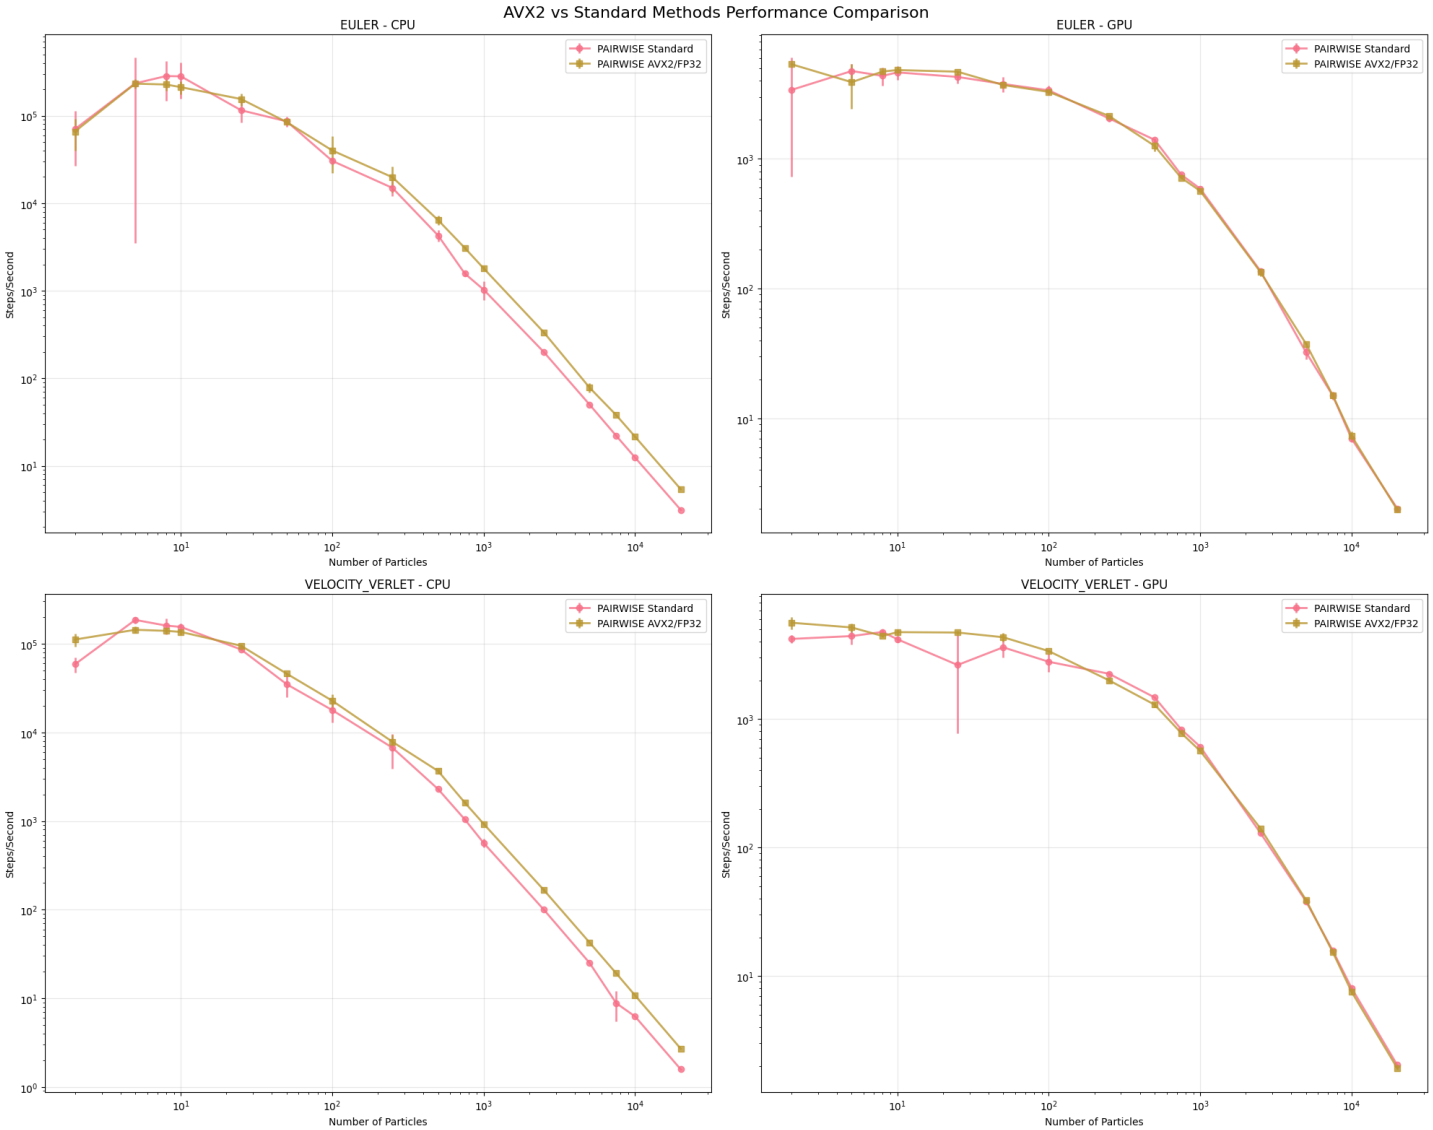

Basic CPU-GPU Comparison Benchmark

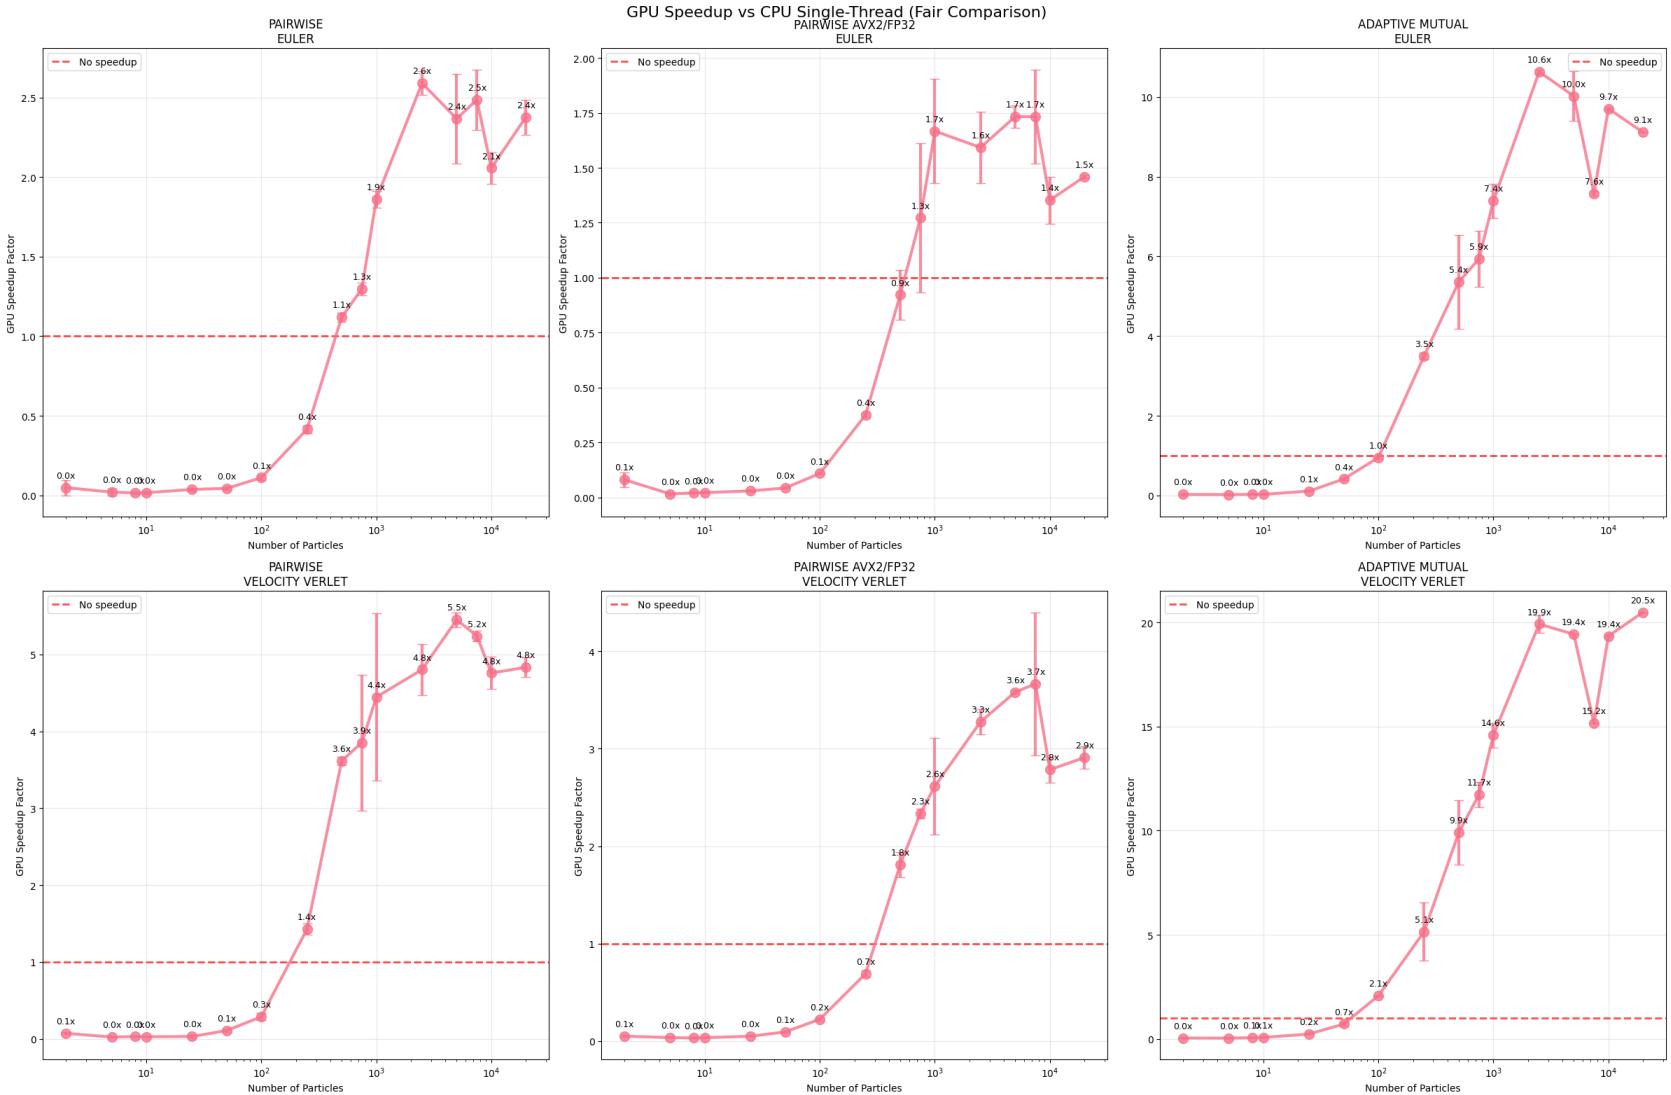

CPU-Methods Comparison Benchmark

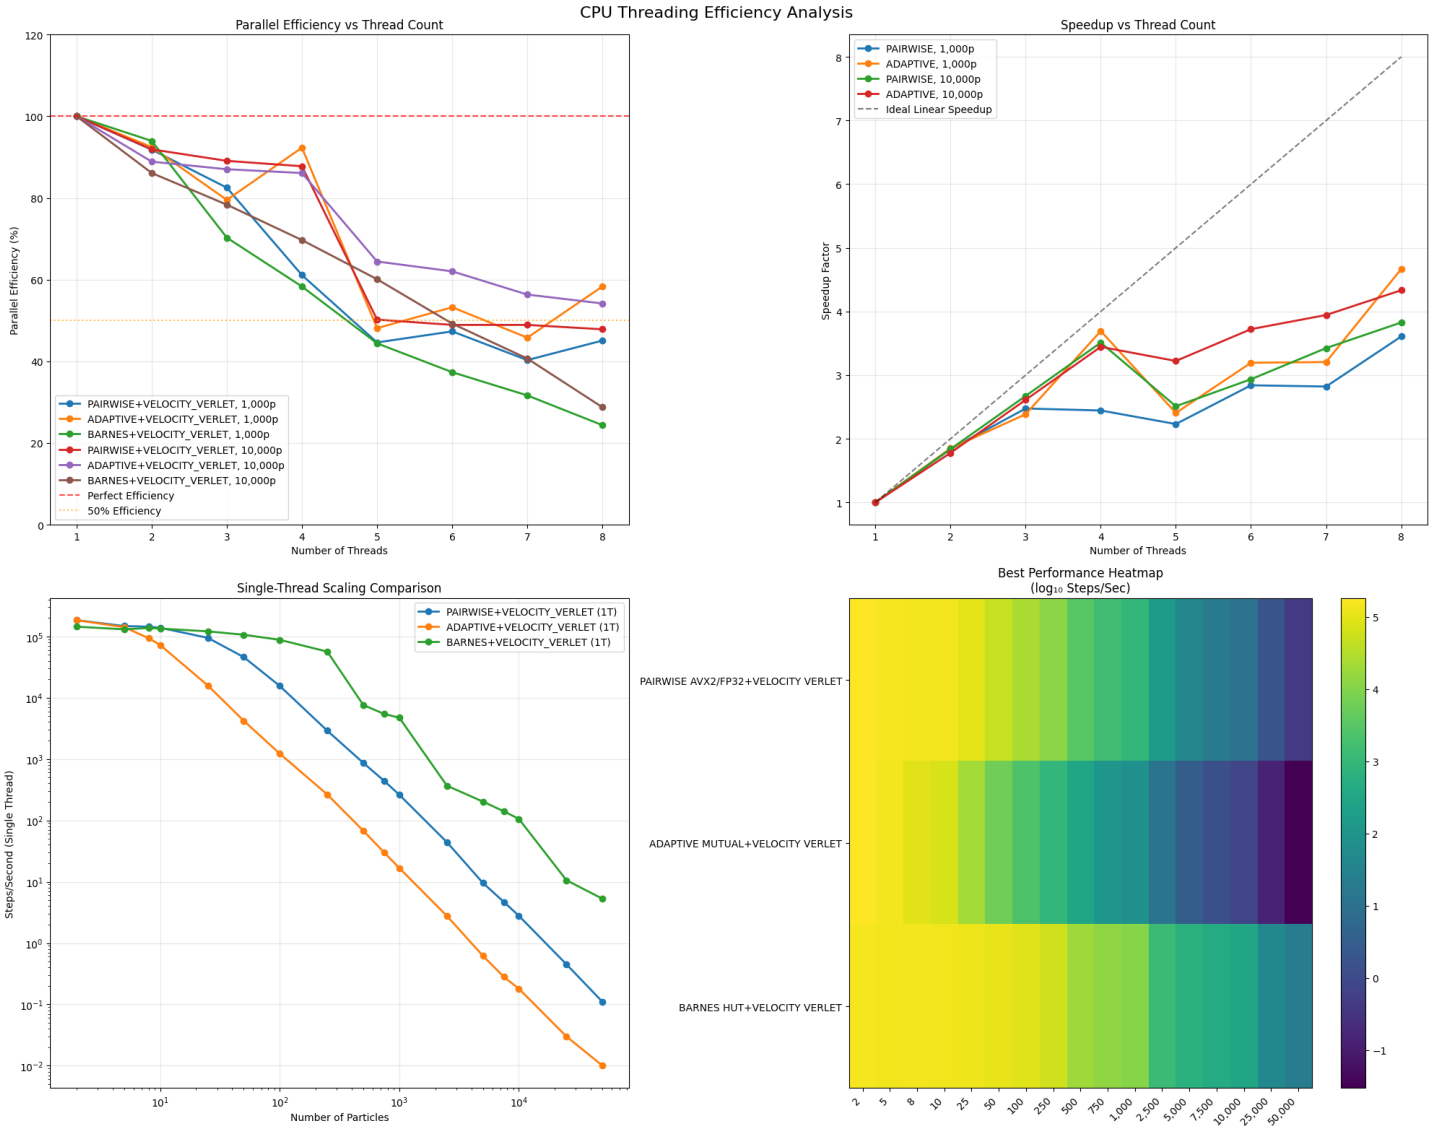

CPU-Methods Comparison Benchmark

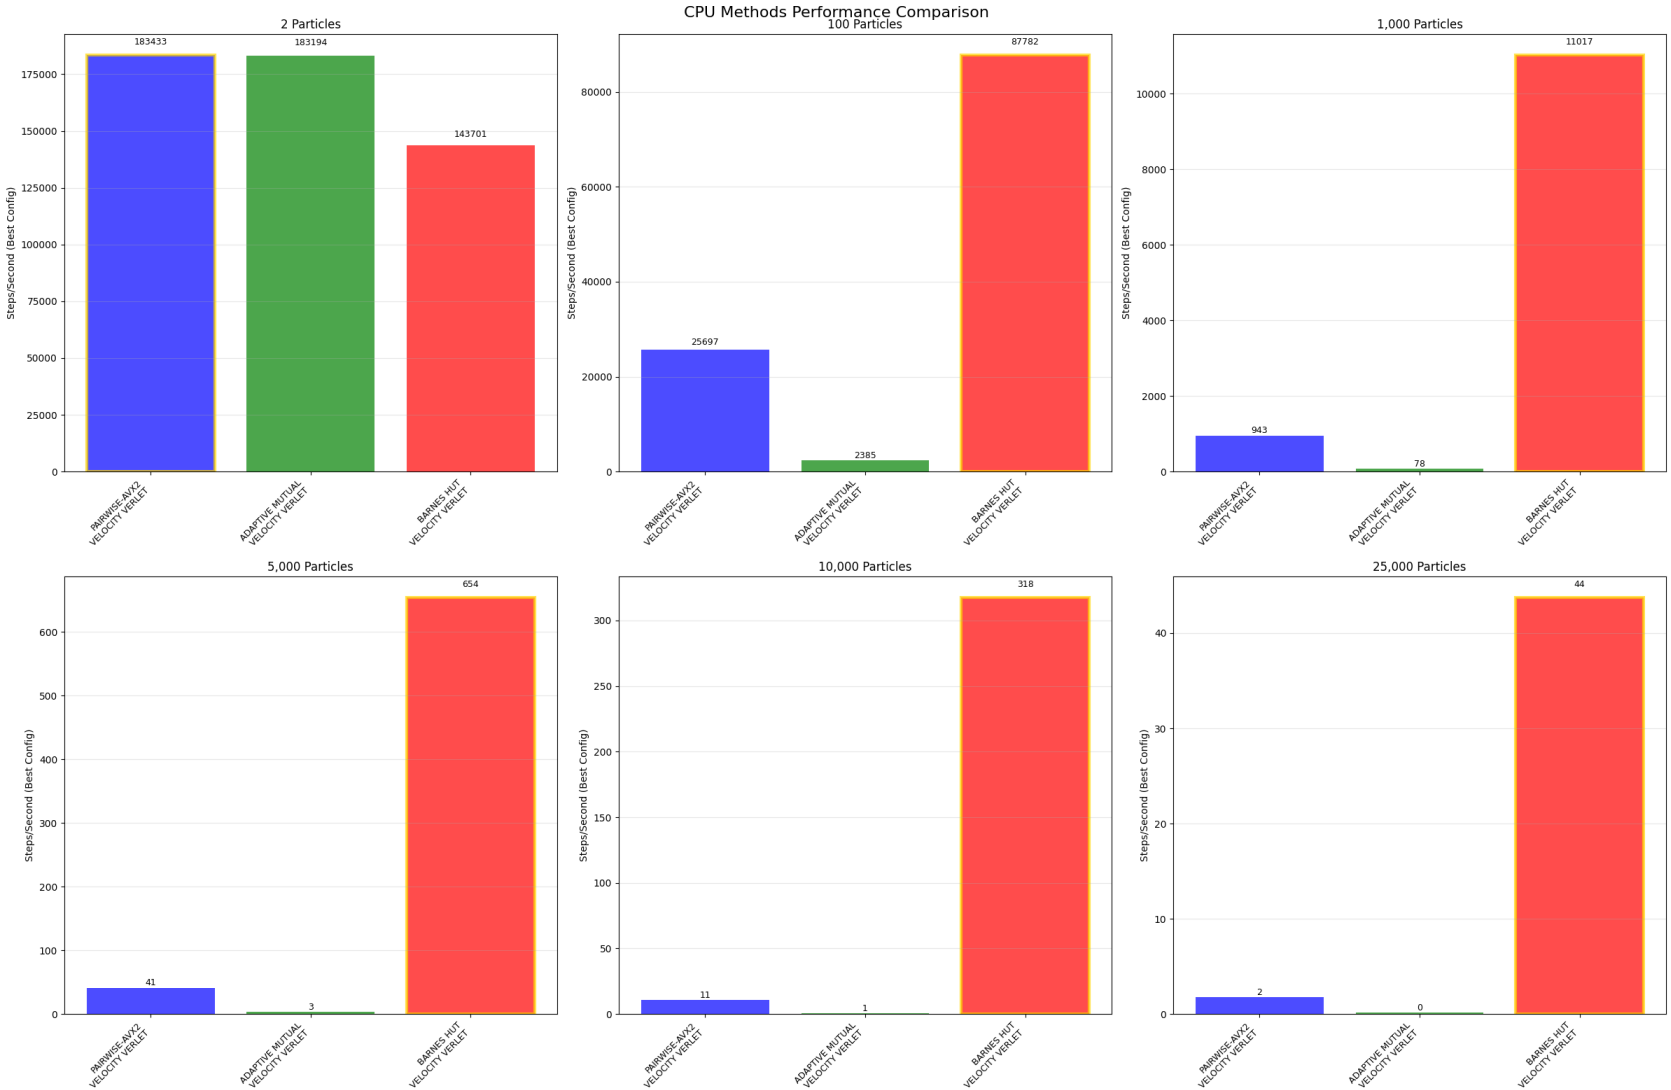

BH vs GPU Comparison Benchmark

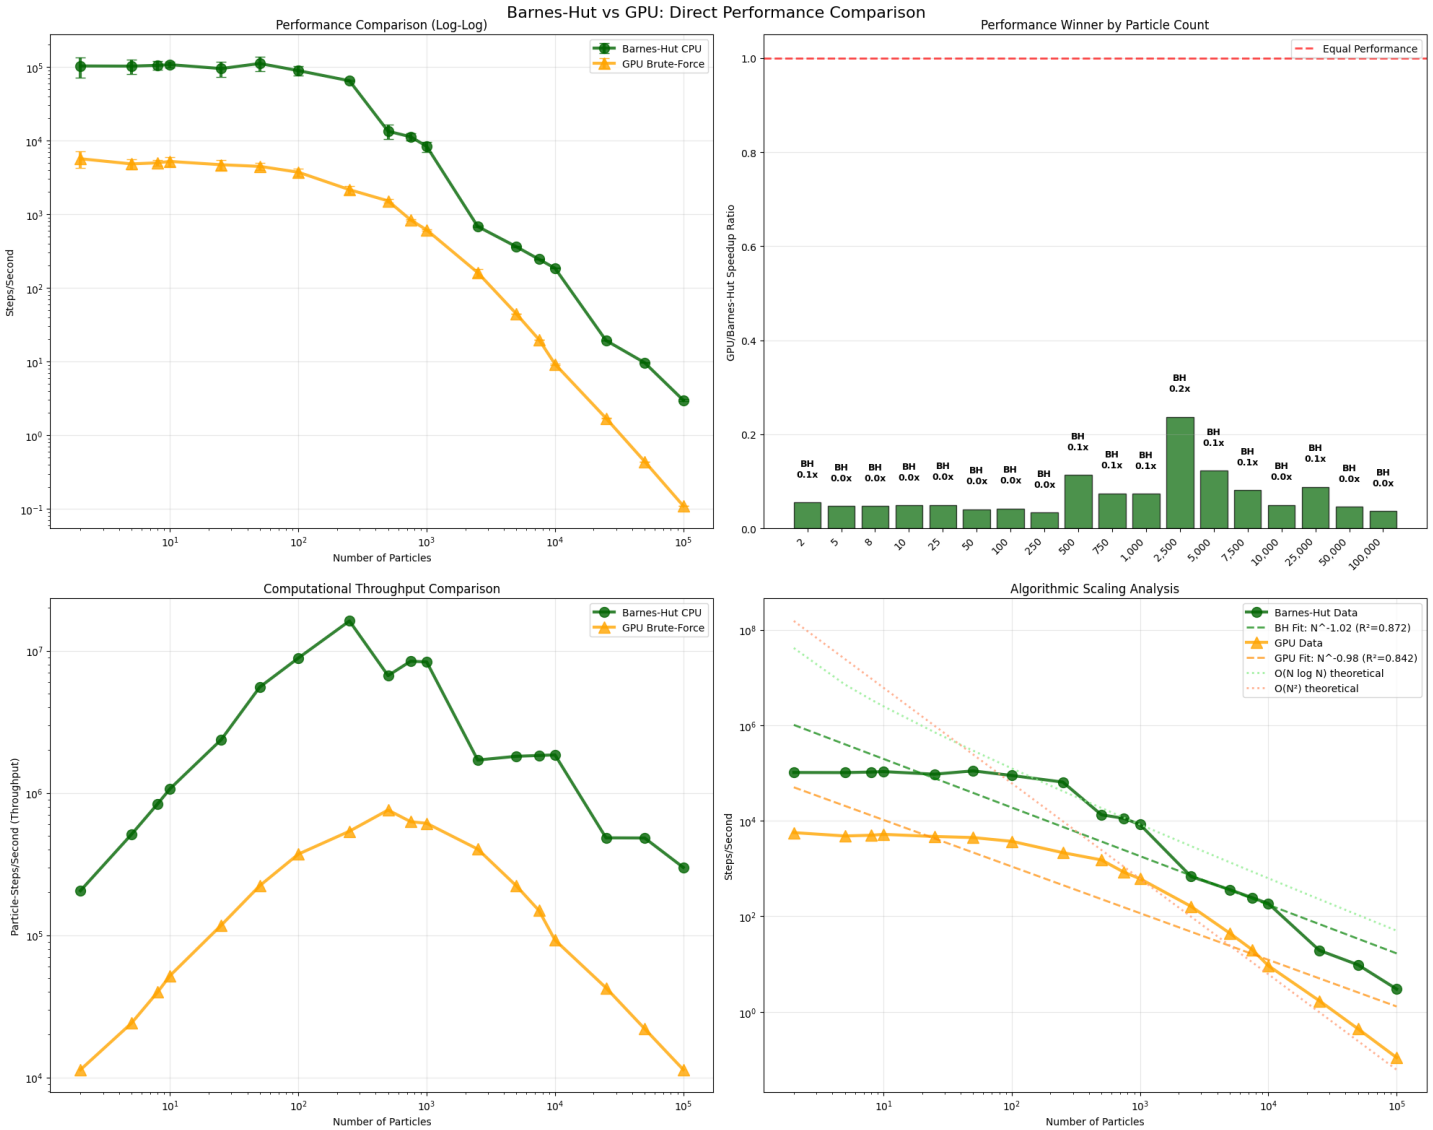

Limitations and Lessons Learned

- Bottlenecks in compute for very large N
- Tradeoffs between accuracy and speed (double in GPU, same memory pattern in CPU).
- Importance of a clear vision at the start of the project
- Need for a robust testing framework
- Make only one main entry point per external library (or as few as possible)
- Future improvements: a running program you can compile without looking at any library requirements

Conclusions

- We skipped many small details (tiling, async memory...) in this presentation, so if you want to know more, please check the GitHub repository
- High-performance, modular N-body simulation achieved
- GPU acceleration enables large-scale experiments
- Framework ready for extension and research

Thank you for your attention!

Questions?

Demo?